

Bayesian Nonparametric Regression Models for Quantifying Complex Interactions in Exposure Mixture Studies

Your R. Name
APRIL DD, 20YY

Submitted to the Department of
Mathematics and Statistics
of Amherst College in partial fulfillment
of the requirements for the degree of
Bachelor of Arts with honors.

ADVISORS:
Advisor F. Name
Your Other Advisor

Abstract

The abstract should be a short summary of your thesis work. A paragraph is usually sufficient here.

Acknowledgments

Use this space to thank those who have helped you in the thesis process (professors, staff, friends, family, etc.). If you had special funding to conduct your thesis work, that should be acknowledged here as well.

Table of Contents

Abstract	i
Acknowledgments	iii
List of Tables	v
List of Figures	vii
Chapter 1: Introduction	1
Chapter 2: Humanistic perspective	5
Chapter 3: Bayesian regression methods	7
3.1 Motivation	7
3.1.1 Interactions from a statistical perspective	7
3.1.2 Mechanistic and public health relevance	9
3.2 Bayesian kernel machine regression	10
3.2.1 Kernel machine regression	11
3.2.2 Connection to mixed models	14
3.2.3 Toy example	15
3.2.4 Variable selection	18
3.2.5 Prior specification	20
3.2.6 The MCMC algorithm	21
3.2.7 Bayesian multiple index model extension	22
3.3 Bayesian semiparametric regression	22

3.3.1	Spline regression	23
3.3.2	Connection to mixed models	23
3.3.3	Toy example	23
3.3.4	Sparsity inducing priors	24
3.3.5	Prior specification	24
3.3.6	The MCMC algorithm	24
3.4	Detecting interactions	24
3.4.1	BKMR	24
3.4.2	BMIM	24
3.4.3	BSR	24
3.4.4	Toy example	25
Chapter 4:	Simulations	27
4.1	Past simulation studies	27
4.2	Methods	27
4.2.1	Models	27
4.2.2	MADRES data	27
4.2.3	Simulating data	28
4.2.4	Model assessment	28
4.3	Results	28
Conclusion	29
Appendix A:	The First Appendix	31
A.1	In the main file ??:	31
A.2	In Chapter ??:	31
Appendix B:	The Second Appendix	33
Corrections	35

References	37
----------------------	----

List of Tables

List of Figures

3.1	Non-linear data with a true distribution (orange) and a fitted linear regression (blue).	15
3.2	A query point of 12.5 and the weights of neighboring observations based on a Gaussian kernel	16
3.3	Fitted kernel machine regression with $\rho = 2$	17
3.4	Fitted kernel machine regression with $\rho = 0.02$ and $\rho = 50$	17
3.5	Linear spline regression	23
3.6	Spline regression	23

Chapter 1 Introduction

Rapid industrial development has created conditions of cumulative chronic toxicity which pose an acute risk to the wellbeing of humans and our living environment. In fact, it has been estimated that human activity releases chemicals at a rate of 220 billion tons per annum (Cribb, 2016). Scholars have recently declared that, at this global rate of chemical release, humanity has now surpassed the safe operating space of the planetary boundary for novel entities (Persson et al., 2022). As a result, exposure to low levels of pollutants has become an inevitable peril of daily life (Naidu et al., 2021; Vineis, 2018). Hence, it is especially timely that we prompt regulatory control of industrial pollution through studies which investigate the health effects of chemical exposures.

For this, we turn to epidemiological studies. The broad field of preventive epidemiology involves the identification of potentially modifiable risk factors that contribute to the burden of disease within human populations. Environmental epidemiology, in particular, considers the effect of environmental exposures — chemical or otherwise. However, studies concerning chemical pollutants in environmental epidemiology have historically focused on elucidating the effect and mechanisms of exposures to a single pollutant. In reality, humans are invariably exposed to numerous complex chemical mixtures which together contribute to the progression of adverse health outcomes. Therefore, risk assessments of single pollutants likely fail to capture the true consequences of these complex exposures (Heys et al., 2016). Assessing mixtures of

chemicals can also have more direct implications for public health interventions. The United States Environmental Protection Agency (U.S. EPA) currently passes regulations for individual pollutants. In practice, though, regulation occurs by controlling the source of pollution, which is responsible for the production of a whole mixture of chemicals with specific joint effects on human health. As a result, the National Academies of Science has advocated for a multipollutant regulatory approach, which is likely to be more protective of human health (Committee on Incorporating 21st Century Science into Risk-Based Evaluations et al., 2017).

There are clear practical motivations for studies that examine the health effects of exposure to co-occurring chemical mixtures, hereafter referred to as exposure mixtures. However, expanding the focus of analysis from one exposure to multiple exposures introduces unique statistical challenges. In addition to a common issue of small effect sizes and small sample sizes present in most exposure analyses, multiple exposure analyses must also contend with high-dimensionality, collinearity, non-linear effects, and non-additive interactions (Yu et al., 2022). In particular, data with numerous pollutants, or predictors, require exponentially greater levels of complexity and time cost in analysis. Collinearity between exposures is common when analyzing pollutants from a single source and can lead to unstable estimates in a generalized linear model if left unaccounted for. Finally, exposures can have both non-linear single effects and non-additive interaction effects, which are difficult to capture unless explicitly specified in the model.

The classic multiple linear regression framework often fails to capture the true effects in this setting. In the past few years, a wide variety of statistical methods have been developed to overcome these challenges (Gibson et al., 2019; Yu et al., 2022), which have been accompanied by a host of comparative simulation studies for general mixture scenarios (e.g., Hoskovec et al., 2021; Lazarevic et al., 2020; Pesenti

et al., 2023). However, to our knowledge, there has yet to be a simulation study which provides conclusive guidance about the ability of these methods to conduct inference on non-additive interactions between exposures, particularly when the nature and effect sizes of these interactions vary.

The goal of this thesis is to fill this gap in the literature by exploring the theory and performance of Bayesian regression techniques for quantifying complex interactions between multiple environmental exposures and related covariates. Specifically, we will compare two recently developed models for estimating the health effects of exposure mixtures: Bayesian Kernel Machine Regression (BKMR, Bobb et al., 2015) and Bayesian Semiparametric Regression (BSR, Antonelli et al., 2020).

In an age where anthropogenic actions have radically reshaped the earth, humanistic inquiry can offer critical insights into how we navigate the hazards of our rapidly changing environment. We begin in Chapter 2 by contextualizing this thesis with a brief overview of cultural and social understandings of toxicity. Chapter 3 explains the motivation for studying interactions and provides background on the theory of Bayesian methods for analyzing exposure mixtures. Chapter 4 assesses the performance of these methods using a simulation study, based on a dataset with information on the relationship between prenatal exposure to heavy metals and gestational weight. Chapter 5 explores an application on X data [TBD]. We conclude with a discussion of the implications of this work for the future study of complex interactions in exposure mixture studies.

Chapter 2 Humanistic perspective

Chapter 3 Bayesian regression methods

3.1 Motivation

We are interested in using Bayesian regression techniques to characterize the nature of complex interactions in exposure mixture studies. We begin in this section by reviewing definitions for what constitutes a complex interaction and their public health and biological relevance.

3.1.1 Interactions from a statistical perspective

First, we define additivity and non-additivity in the traditional statistical paradigm (Siemiatycki & Thomas, 1981). Suppose we have two variables x_1 and x_2 , and we want to consider their effect on some outcome of interest. If specifying [effect due to x_1 and x_2] = [effect due to x_1] + [effect due to x_2] can adequately capture this relationship, then we say that x_1 and x_2 each have an **additive effect** on the outcome and that there is no interaction between them. On the other hand, if there is variability in the outcome that can be captured by an additional term equal to some function of x_1 and x_2 , we say that there is a **non-additive interaction** between x_1 and x_2 . In this case, [effect due to x_1 and x_2] = [effect due to x_1] + [effect due to x_2] + [effect due to $f(x_1, x_2)$], where f is non-zero for some values of x_1 and x_2 .

For our sake, we consider any non-additive interaction to be complex. Complex interactions can be difficult to detect. To see why, let us consider running a linear

regression for Y on x_1 and x_2 . The regression would be defined as

$$\hat{Y} = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_{12} f(x_1, x_2),$$

where the β 's represent the effect sizes. We can see that the form of the interaction must be explicitly specified in the formulation of the model. Most commonly, a multiplicative interaction is assessed, where $f(x_1, x_2) = x_1 * x_2$. However, a non-additive interaction can take on many different forms, the true nature of which is difficult to determine analytically.

We used a two-predictor case above, but interactions can also exist between two or more variables (i.e., two-way by $f(x_1, x_2)$, three-way by $f(x_1, x_2, x_3)$, etc.). And if we wanted to assess all possible interactions, the number to consider quickly becomes intractable in high-dimensional settings. For instance, consider modelling 10 predictors in the above linear regression setting. In order to be assessed, each interaction must be explicitly specified as a new term in the model. Even if we only considered one form for the interaction, including all possible two-way interactions would involve adding $\binom{10}{2} = 45$ additional terms to the model, and all possible three-way interactions would add $\binom{10}{3} = 120$ additional terms.

It is important also to acknowledge, here, that there is a limit to how many variables can be included in an interaction before it becomes incomprehensible to most humans. For instance, Halford, Baker, McCredden, & Bain (2005) suggest that there is a steep decline in interpretability from three- to four-way interactions, and that five-way interactions are only interpreted correctly at chance level (Halford et al., 2005). Hence, for practical purposes, we will limit our exploration to two- and three-way interactions.

3.1.2 Mechanistic and public health relevance

Thus far, we have discussed interactions within a statistical paradigm. However, in addition to being an interesting estimation challenge, non-additive interactions are also relevant in exposure mixture studies from both a mechanistic and public health point of view.

From a mechanistic perspective, a non-additive statistical interaction between two chemical exposures suggests that these molecules may be functionally interacting with each other. Theoretical models propose that such interactions can be classified as either synergistic or antagonistic (Heys, Shore, Pereira, Jones, & Martin, 2016; Plackett & Hewlett, 1952). In a synergistic interaction, the joint effects of a mixture exceed the independent effects of each component. This usually occurs if a chemical induces an enzyme involved with the activation of a second chemical or if a chemical inhibits an enzyme that would have otherwise degraded a second chemical. For example, various synergistic interactions between heavy metals have been documented in the literature (e.g., Paithankar, Saini, Dwivedi, Sharma, & Chowdhuri, 2021; Wang, Mukherjee, & Park, 2018).

On the other hand, in an antagonistic interaction, the joint effects of a mixture are less than their independent effects. This can occur either through competition at the target site of an enzyme or through direct chemical reactions with each other. In general, synergistic interactions are more concerning in risk assessments, as they lead to underestimation of the true toxicity of a mixture.

It should be noted, though, that while statistical interactions may provide some insight into how exposure mixtures are related to health, they cannot confirm their underlying biology (VanderWeele & Knol, 2014). If the goal is to assess a meaningful biological interaction, then the discovery of a statistical interaction should be followed

up by a functional study.

Now, from a public health perspective, we might be interested in how exposure mixtures interact with other covariates, or, in other words, how social and health factors might mediate the relationship between a health outcome and chemical exposures (VanderWeele & Knol, 2014). In our case, we can include these additional covariates in the exposure mixture model, where, statistically, they would contribute to the model in the same manner as another chemical exposure: a predictor. A statistical interaction in our model between a covariate and an exposure would indicate that the *magnitude* of the effect of reducing the level of an exposure might differ across various levels of the covariate. This finding would be relevant to public health policy makers, as the potential benefit of regulating a pollutant might differ across groups. For instance, it has been suggested that nutritional intake may modify susceptibility to chemical exposures (e.g., Kannan, Misra, Dvonch, & Krishnakumar, 2006; Kordas, Lönnerdal, & Stoltzfus, 2007)

In many cases, we might assess a covariate related to health inequity, such as socioeconomic status. We provide a cautionary comment, here, that an interaction term should not be the *sole* measure used to measure a health disparity (Ward et al., 2019). In this case, we should first consider the independent, additive association between the covariate and levels of exposure or rates of a health outcome, in order to contextualize the meaning of a potential interaction term.

3.2 Bayesian kernel machine regression

In this section, we introduce the theory of BKMR. First, we define the notation that we will be using for kernel machine regression:

- \mathbf{x}_m is a predictor variable in the predictor matrix \mathbf{X} with $m = 1, \dots, M$, mea-

measuring exposure variables or covariates in this case

- \mathbf{x}_i is a vector of values for a single observation in \mathbf{X} with $i = 1, \dots, n$
- x_{im} is an observation of \mathbf{x}_m
- Y_i is an observation of \mathbf{Y} , measuring the health outcome in this case
- $h(\cdot)$ is the flexible function relating \mathbf{x} to \mathbf{Y}
- k is the kernel function, the Gaussian in this case
- \mathbf{K} is the $n \times n$ kernel matrix, with (i, j) th element $k(\mathbf{x}_i, \mathbf{x}_j)$
- ρ is the parameter inside the kernel function which controls smoothness
- τ is the parameter multiplied by the kernel matrix to relate \mathbf{K} to h
- $\epsilon_i \stackrel{\text{iid}}{\sim} N(0, \sigma^2)$ are the residuals of the response

And, we define the notation that we will be using specific to BKMR:

- $r_m = 1/\rho_m$ is an augmented variable in \mathbf{r} in the kernel matrix, controlling smoothness
- δ_m is an indicator variable in $\boldsymbol{\delta}$ which represents inclusion in the model
- \mathcal{S}_g is a group of partitioned predictors with $g = 1, \dots, G$
- $\{\delta_m | \mathbf{x}_m \in \mathcal{S}_g\}$ is an indicator variable in $\boldsymbol{\delta}_{\mathcal{S}_g}$ which represents inclusion of a parameter in group g in the model
- π is the prior probability of inclusion of a predictor in the model
- $\lambda \equiv \tau\sigma^{-2}$ is a convenient way to define the prior on τ

3.2.1 Kernel machine regression

We begin by introducing kernel machine regression, with attention to its specific implementation in BKMR. First proposed by Nadaraya (1964) and Watson (1964), kernel machine regression is a nonparametric regression technique that can be used to capture non-linear effects and non-additive interactions. In this introduction, we follow the presentation of kernel machine regression provided by Bobb et al. (2015).

To contextualize this method, we start at the typical linear regression setting,

$$Y_i = \mathbf{x}_i^\top \boldsymbol{\beta} + \epsilon_i,$$

where Y_i measures a health outcome at a given point, $\mathbf{x}_i = [x_1, \dots, x_M]^\top$ is a vector of M exposures or covariates (hereafter referred to as predictors), $\boldsymbol{\beta}$ is a vector of weights, and ϵ_i is a random variable from $\epsilon \stackrel{\text{iid}}{\sim} N(0, \sigma^2)$. We can see that this function assumes that there is a linear relationship between the exposure and the response, and that the combined effects of multiple exposures are additive.

Kernel machine regression defines this relationship using a flexible function $h : \mathbb{R}^M \rightarrow \mathbb{R}$, where

$$Y_i = h(\mathbf{x}_i) + \epsilon_i.$$

Here, $h(\cdot)$ is represented by the function $k(\cdot, \cdot)$, a kernel. The kernel controls the covariance, or the similarity, between values of $h(\mathbf{x})$ and as such ensures that points near each other on the prediction surface will have similar values — or, in other words, that the prediction surface will be smooth. In the case of kernel machine regression, we define a positive definite kernel where $k : \mathbb{R}^M \times \mathbb{R}^M \rightarrow \mathbb{R}$.

There are many choices of functions for k . BKMR uses the Gaussian kernel, also known as the radial basis function or, sometimes, the squared exponential kernel. The Gaussian kernel is defined as

$$k(\mathbf{x}, \mathbf{x}') = \exp\left\{-\frac{\|\mathbf{x} - \mathbf{x}'\|^2}{\rho}\right\},$$

where $\|\mathbf{x} - \mathbf{x}'\|^2 = \sum_{m=1}^M (x_m - x'_m)^2$ for a set of predictors values \mathbf{x} and the predictor values of a second subject \mathbf{x}' , and ρ is a tuning parameter that controls the relationship between the correlation between two points and their distance. Greater values of ρ will

enforce more dependence between points and make the resulting function smoother. h is related to k by a multiplicative constant τ , a tuning parameter which controls the vertical scale of h .

Now that we have defined h and k , we can think about how to characterize the relationship between our response and predictors. Kernel machine regression is a nonparametric technique because it does not specify a functional form for this relationship. Hence, we will think about estimating the response at a particular query point. Operationally, Müller (1987) demonstrated that kernel machine regression uses a weighted average of all the observations in the dataset to estimate the response, defined as

$$\bar{Y} = \frac{\sum_{i=1}^n w_i Y_i}{\sum_{i=1}^n w_i},$$

with some set of weights $\{w_i\}_{i=1}^n$. Intuitively, we want to weight the observations that are closer to the query point more heavily. Using the Gaussian kernel as a weight allows us to achieve this. Replacing the weight with the Gaussian kernel, we get

$$\bar{Y} = \frac{\sum_{i=1}^n k(\mathbf{x}, \mathbf{x}_i) Y_i}{\sum_{i=1}^n k(\mathbf{x}, \mathbf{x}_i)}.$$

As we move through the predictor space, we can think of the prediction as a continuous moving average of local points in the dataset. The correlation between two values of h is defined as

$$\text{cor}(h_i, h_j) = \exp\left\{-\frac{\|x_{im} - x_{jm}\|}{\rho}\right\},$$

which allows us to see that values of h near each other will have a higher correlation and thus similar values. This is also why the resulting function is smooth.

3.2.2 Connection to mixed models

It is useful to make connections between this definition of kernel machine regression and mixed models. Liu, Lin, & Ghosh (2007) demonstrated this by representing $h(\mathbf{x})$ as following a Gaussian process probability distribution,

$$h(\mathbf{x}) \sim \mathcal{GP}(\mathbf{0}, \tau k(\mathbf{x}, \mathbf{x}')),$$

with covariance function k , where \mathbf{x} is a vector of the predictor values, and \mathbf{x}' contains the predictor values of another subject. A Gaussian process is a collection of random variables, of which any finite number follow a multivariate normal distribution (Schulz, Speekenbrink, & Krause, 2018). Here, we assume that the expected value of the h function with input \mathbf{x} is $\mathbf{0}$. We use k for the covariance function, which represents the dependence between the function values with inputs \mathbf{x} and \mathbf{x}' : $k(\mathbf{x}, \mathbf{x}') = \mathbb{E}[(h(\mathbf{x}) - \mathbf{0})(h(\mathbf{x}') - \mathbf{0})]$.

Now, we can represent h as a collection of variables from a Gaussian process. h follows a multivariate normal distribution,

$$h(\mathbf{x}) \sim N(\mathbf{0}, \tau \mathbf{K}),$$

where $h(\mathbf{x}) = [h(\mathbf{x}_1), h(\mathbf{x}_2), \dots, h(\mathbf{x}_n)]^\top$ and \mathbf{K} is the kernel matrix. The kernel matrix is an $n \times n$ matrix with (i, j) th element $k(\mathbf{x}_i, \mathbf{x}_j)$. Now, returning back to the regression view, we can think of each Y_i as following the distribution,

$$Y_i \stackrel{\text{ind}}{\sim} N(h(\mathbf{x}_i), \sigma^2) \text{ for } i = 1, \dots, n,$$

where σ^2 comes from the variance of the residuals. Here, h can be interpreted as a random effect.

3.2.3 Toy example

In the following section, we illustrate kernel machine regression with a toy example.

Consider the following case where we want to model the relationship between a single predictor and a response variable. Suppose the true relationship between x and Y is defined $Y = e^{\frac{x}{10}} + 2 \sin(\frac{x}{2})$.




Figure 3.1: Non-linear data with a true distribution (orange) and a fitted linear regression (blue).

Figure 3.1 illustrates the shape of our simulated non-linear data and the fit proposed by a simple linear regression. We can observe that the linear regression fails to capture the true non-linear relationship. In this case, this would lead to an underestimation of the true association between x and Y . Now, we will try to capture this relationship using kernel machine regression.

To visualize how kernel machine regression works as a moving weighted average, we can consider a query point of 12.5.




Figure 3.2: A query point of 12.5 and the weights of neighboring observations based on a Gaussian kernel

Figure 3.2 identifies the query point and assigns corresponding weights to the neighboring points based on a normal distribution, which shares the same density as the Gaussian kernel. In this case, we will specify $\rho = 2$, which is synonymous with assigning the weights using a normal distribution with $\sigma^2 = 1$. We can see how an appropriate estimate for $h(12.5)$ can be obtained by taking a weighted average of the x 's, with those observations nearby weighted the most heavily.

Now, we fit a kernel machine regression on this data with $\rho = 2$ using the `stats` package in R (R Core Team, 2013).




Figure 3.3: Fitted kernel machine regression with $\rho = 2$.

We can see in figure 3.3 that kernel machine regression captures the complex non-linear relationship between Y and x and closely follows the true distribution. We do note, though, that the estimation is less precise at the tails, where there is less information provided by local observations. We can also use this example to consider the effect of various values of ρ on the smoothness of the h function.




Figure 3.4: Fitted kernel machine regression with $\rho = 0.02$ and $\rho = 50$.

Figure 3.4 demonstrates the effect of relatively smaller and larger values of ρ on

h . Decreasing the value of ρ allows kernel machine regression to overfit to the noise in the data by relaxing the dependence of neighboring values of h to each other. On the other hand, increasing the value of ρ enforces more dependence in h and as such results in an underfit estimation. Hence, the choice of ρ has a strong effect on the performance of kernel machine regression.

3.2.4 Variable selection

Now that we have defined kernel machine regression, we can extend it to the Bayesian paradigm. bobb_bayesian_2015 showed that the Bayesian approach can outperform frequentist kernel machine regression because simultaneous variable selection and estimation can better capture the exposure-response relationship. In this section, we discuss the two methods for Bayesian variable selection in BKMR: hierarchical variable selection and component-wise variable selection (Bobb et al., 2015).

Component-wise selection follows the same framework as variable selection in a typical Bayesian multiple regression except, instead of augmenting each predictor, we augment the kernel function as

$$k(\mathbf{x}, \mathbf{x}' | \mathbf{r}) = \exp\left\{-\sum_{m=1}^M r_m (x_m - x'_m)^2\right\},$$

where $\mathbf{r} = [r_1, \dots, r_M]^\top$. We define $r_m = 1/\rho_m$, the inverse of the tuning parameter ρ_m for each \mathbf{x}_m , as we want r_m to be able to take on the value of 0 when a variable is removed during the selection process. We now define the kernel matrix $\mathbf{K}_{\mathbf{x}, \mathbf{r}}$ as the $n \times n$ matrix with (i, j) th element $k(\mathbf{x}_i, \mathbf{x}'_j | \mathbf{r})$. To allow r_m to equal 0 with non-zero probability, we first define an indicator variable determining whether or not a predictor is included in the variable, which is distributed as

$$\delta_m \sim \text{Bernoulli}(\pi),$$

where π is the prior probability of inclusion. Now, we can assume a “slab-and-spike” prior on r_m , distributed as

$$r_m | \delta_m \sim \delta_m f(r_m) + (1 - \delta_m) P_0,$$

where $f(\cdot)$ is some pdf with support \mathbb{R}^+ and P_0 denotes the density with point mass at 0.

While this process of component-wise variable selection works well in a typical multiple regression setting, it can lead to unreliable estimates in situations where the predictors are highly correlated with each other, which is common in exposure mixture studies. In this case, the correlated components contribute similar information to the model, and component-wise variable selection is not able to distinguish which predictor is important. BKMR deals with this problem by introducing hierarchical variable selection.

Hierarchical variable selection involves partitioning the predictors $\mathbf{x}_1, \dots, \mathbf{x}_M$ a priori into groups \mathcal{S}_g with $g = 1, \dots, G$. These groups should be selected with the aim of keeping within-group correlation high and between-group correlation low. The indicators from $r_m | \delta_m$ are now distributed as

$$\boldsymbol{\delta}_{\mathcal{S}_g} | \omega_g \sim \text{Multinomial}(\omega_g, \boldsymbol{\pi}_{\mathcal{S}_g}), g = 1, \dots, G, \omega_g \sim \text{Bernoulli}(\pi),$$

where $\boldsymbol{\delta}_{\mathcal{S}_g} = \{\delta_m | \mathbf{x}_m \in \mathcal{S}_g\}$ and $\boldsymbol{\pi}_{\mathcal{S}_g}$ are vectors of indicator variables and prior probabilities, respectively, of a predictor \mathbf{x}_m in group \mathcal{S}_g entering the model. By this approach, at most one predictor in each group is allowed to enter the model.

While hierarchical variable selection resolves the issue of multicollinearity, it re-

quires specifying subgroups of predictors a priori and assumes that one predictor in each group can capture the information of the rest. Hence, care should be taken to justify the partitioning of predictors when taking this approach.

Note also that the posterior means of δ_m generated from these variable selection procedures represent the posterior probability of inclusion of \mathbf{x}_m . We can interpret these posterior “inclusion probabilities” as measures of the relative importance of each predictor. These measures can be used to understand the contribution of each exposure or covariate to the health outcome of interest in the model.

3.2.5 Prior specification

In this section, we specify the prior distributions and parameters used by the BKMR algorithm (Bobb et al., 2015).

BKMR, by default, assumes $\rho_m = 1/r_m \sim \text{Unif}(a_r, b_r)$, a flat prior between a_r and b_r for which the default values are 0 and 100, respectively (Bobb, 2017a). This defines the prior probability of ρ as equally distributed across any value from 0 to 100. This inverse of this prior corresponds to the slab component of the “slab-and-spike” prior, where $r_m|\delta_m \sim \delta_m \text{Unif}^{-1}(a_r, b_r) + (1 - \delta_m)P_0$. As a flat prior, this distribution should be chosen when we have no prior knowledge about the smoothness of the exposure-response function, with hyperparameters a_r and b_r selected to represent the range of values we expect ρ to potentially span.

We have seen that kernel machine regression responds strongly to different values of $r_m = 1/\rho$, and, accordingly, the model fit of BKMR is sensitive to their prior distribution. In general, the posterior inclusion probabilities generated from the variable selection procedure are particularly sensitive to this prior, though their relative importance tends to remain stable (Bobb, Claus Henn, Valeri, & Coull, 2018). As such, the BKMR algorithm also offers the options to define uniform and gamma priors for

the $r_m = 1/\rho$.

Moreover, BKMR assumes that the prior probability of including a predictor (δ_m) or group of predictors (ω_g) in the model is distributed $\pi \sim \text{Beta}(a_\pi, b_\pi)$. The default hyperparameters are $a_\pi = b_\pi = 1$, which represent a flat, uninformative prior between 0 and 1. When the hierarchical selection approach is applied, equal values for $\boldsymbol{\pi}_{\mathcal{S}_g}$, the probabilities of inclusion for each component in group \mathcal{S}_g , are assumed.

Finally, BKMR assumes that the inverse of the variance of the residuals is distributed $\sigma^{-2} \sim \text{Gamma}(a_\sigma, b_\sigma)$, with default values of $a_\sigma = b_\sigma = 0.001$, and that the vertical scale of h is parameterized by $\lambda \equiv \tau\sigma^{-2} \sim \text{Gamma}(a_\lambda, b_\lambda)$, with default values of a_λ, b_λ such that the mean and variance of λ are both equal to 10.

3.2.6 The MCMC algorithm

Briefly, we discuss the algorithm used in the BKMR package (Bobb et al., 2015, 2018), with commentary on its implications for the model fitting process.

BKMR uses a Markov chain Monte Carlo (MCMC) algorithm with a mix of Gibbs and Metropolis-Hastings samplers to estimate the posterior distributions. In particular, a Gibbs step is used to update the distribution of σ^2 while a Metropolis-Hastings step is used to update the distribution of λ . For component-wise and hierarchical variable selection, $(\mathbf{r}, \boldsymbol{\delta}, \boldsymbol{\omega})$ are sampled jointly using a Metropolis-Hastings sampling scheme.

While each distribution generated by the Gibbs step is always accepted, the distributions for λ and r_m generated by the Metropolis-Hastings steps are accepted based on an acceptance rate (Wagaman & Dobrow, 2021). The standard deviation of the proposal distribution controls the acceptance rate and as such acts as a tuning parameter (Bobb, 2017b). In general, increasing the standard deviation leads to lower acceptance rates. Acceptance rates that are too low lead to slower convergence, but

rates that are too high can cause convergence to a non-optimal distribution.

A major computational limitation of BKMR is that at each iteration of the MCMC algorithm, the $n \times n$ augmented kernel matrix $\mathbf{K}_{\mathbf{Z},\mathbf{r}}$ must be inverted multiple times. BKMR can employ a Gaussian predictive process which involves specifying a set of l points, or “knots,” that are a subset of the predictor space. The vector of predictors can be approximated by projection on this lower dimensional space, which allows the algorithm to perform inversions on an $l \times l$ matrix. A general suggestion is to use this approach to speed up the algorithm when n is large and to specify $l \approx n/10$ (Bellavia, 2021).

3.2.7 Bayesian multiple index model extension

We briefly introduce an extension. . . *decide whether to keep this section later*

A recent proposal by McGee, Wilson, Webster, & Coull (2023)

- incorporate weighted quantile sums regression + g computation
- cite alex’s paper w/ bayesian approach
- allows variables to contribute through a weighted quantile index.

provides a promising extension to BKMR that resolves the issue with hierarchical variable selection, so we might consider it in our simulations

3.3 Bayesian semiparametric regression

differences w/ bkmr

- makes distributional assumptions about the dataset
- kernel regression is computationally intensive w/ large datasets but can handle many predictors
- bsr highly dependent on the choice of function

3.3.1 Spline regression

3.3.2 Connection to mixed models

3.3.3 Toy example




Figure 3.5: Linear spline regression




Figure 3.6: Spline regression

3.3.4 Sparsity inducing priors

3.3.5 Prior specification

3.3.6 The MCMC algorithm

3.4 Detecting interactions

Discuss options for detecting interactions. Briefly describe MLR

3.4.1 BKMR

Since the flexible h function in kernel machine regression allows us to forgo any assumptions about the nature of the relationship between the health outcome and predictors, BKMR can potentially capture complex interactions between predictors. The challenge with using BKMR to do this, however, is that there is no formal framework for conducting inference on the presence of interactions.

And, if we specify hierarchical variable selection, then

How to detect interactions:

- can visualize relationship at various quantiles of other predictor
- can conduct inference on inter-quantile difference — has yet to be a formal assessment of how this approach performs in various settings

3.4.2 BMIM

- note how interactions work w/in groups and b/t groups

3.4.3 BSR

BSR can conduct inference

3.4.4 Toy example

Potentially provide toy example

Chapter 4 Simulations

4.1 Past simulation studies

Lit review of past sim studies which include bayesian regression approaches + others, note how they haven't focused on interactions

4.2 Methods

What are our aims?

4.2.1 Models

Models compared, specify the parameters for each (justify them!)

- MLR
- MLR with known form of interactions specified
- BKMR with component-wise
- BKMR with hierarchical
- BMIM? extension of BKMR with hierarchical
- BSR

4.2.2 MADRES data

- cite original study design paper

- cite Howe et al. 2021, which applied BKMR and BSR to this data, what did they find?

4.2.3 Simulating data

- use original predictor values in MADRES data. allows us to explore these methods on real-world data
- simulate outcome using a formula w/ different types of interaction
- change effect size (three levels?)
- change nature of interaction (mathematical formulation, two- or three-way, b/t just exposures or b/t exposures and covariates)

4.2.4 Model assessment

- how many times is interaction picked up?

4.3 Results

- figures + tables w/ model performance

Conclusion

If we don't want the conclusion to have a chapter number next to it, we can add the `{-}` attribute.

More info

And here's some other random info: the first paragraph after a chapter title or section head *shouldn't be* indented, because indents are to tell the reader that you're starting a new paragraph. Since that's obvious after a chapter or section title, proper typesetting doesn't add an indent there.

Appendix A The First Appendix

This first appendix includes all of the R chunks of code that were hidden throughout the document (using the `include = FALSE` chunk tag) to help with readability and/or setup.

A.1 In the main file ??:

A.2 In Chapter ??:

Appendix B The Second Appendix

R code

Corrections

A list of corrections after submission to department.

Corrections may be made to the body of the thesis, but every such correction will be acknowledged in a list under the heading “Corrections,” along with the statement “When originally submitted, this honors thesis contained some errors which have been corrected in the current version. Here is a list of the errors that were corrected.” This list will be given on a sheet or sheets to be appended to the thesis. Corrections to spelling, grammar, or typography may be acknowledged by a general statement such as “30 spellings were corrected in various places in the thesis, and the notation for definite integral was changed in approximately 10 places.” However, any correction that affects the meaning of a sentence or paragraph should be described in careful detail. The files `samplethesis.tex` and `samplethesis.pdf` show what the “Corrections” section should look like. Questions about what should appear in the “Corrections” should be directed to the Chair.

References

- Bellavia, A. (2021). *Statistical Methods for Environmental Mixtures*. Retrieved from <https://bookdown.org/andreabellavia/mixtures/preface.html>
- Bobb, J. F. (2017a, March). Introduction to Bayesian kernel machine regression and the bkmr R package. Retrieved from <https://jenfb.github.io/bkmr/overview.html>
- Bobb, J. F. (2017b, December). Example using the bkmr R package with simulated data from the NIEHS mixtures workshop. Retrieved from https://jenfb.github.io/bkmr/SimData1.html#1_load_packages_and_download_data
- Bobb, J. F., Claus Henn, B., Valeri, L., & Coull, B. A. (2018). Statistical software for analyzing the health effects of multiple concurrent exposures via Bayesian kernel machine regression. *Environmental Health*, 17(1), 67. <http://doi.org/10.1186/s12940-018-0413-y>
- Bobb, J. F., Valeri, L., Claus Henn, B., Christiani, D. C., Wright, R. O., Mazumdar, M., ... Coull, B. A. (2015). Bayesian kernel machine regression for estimating the health effects of multi-pollutant mixtures. *Biostatistics*, 16(3), 493–508. <http://doi.org/10.1093/biostatistics/kxu058>
- Halford, G. S., Baker, R., McCredde, J. E., & Bain, J. D. (2005). How Many Variables Can Humans Process? *Psychological Science*, 16(1), 70–76.

<http://doi.org/10.1111/j.0956-7976.2005.00782.x>

- Heys, K., Shore, R., Pereira, M., Jones, K., & Martin, F. (2016). Risk assessment of environmental mixture effects. *RSC Advances*, 6(53), 47844–47857. <http://doi.org/10.1039/C6RA05406D>
- Kannan, S., Misra, D. P., Dvornch, J. T., & Krishnakumar, A. (2006). Exposures to Airborne Particulate Matter and Adverse Perinatal Outcomes: A Biologically Plausible Mechanistic Framework for Exploring Potential Effect Modification by Nutrition. *Environmental Health Perspectives*, 114(11), 1636–1642. <http://doi.org/10.1289/ehp.9081>
- Kordas, K., Lönnerdal, B., & Stoltzfus, R. J. (2007). Interactions between nutrition and environmental exposures: Effects on health outcomes in women and children. *The Journal of Nutrition*, 137(12), 2794–2797. <http://doi.org/10.1093/jn/137.12.2794>
- Liu, D., Lin, X., & Ghosh, D. (2007). Semiparametric Regression of Multidimensional Genetic Pathway Data: Least-Squares Kernel Machines and Linear Mixed Models. *Biometrics*, 63(4), 1079–1088. <http://doi.org/10.1111/j.1541-0420.2007.00799.x>
- McGee, G., Wilson, A., Webster, T. F., & Coull, B. A. (2023). Bayesian multiple index models for environmental mixtures. *Biometrics*, 79(1), 462–474. <http://doi.org/10.1111/biom.13569>
- Müller, H.-G. (1987). Weighted Local Regression and Kernel Methods for Nonparametric Curve Fitting. *Journal of the American Statistical Association*, 82(397), 231–238. <http://doi.org/10.1080/01621459.1987.10478425>
- Nadaraya, E. A. (1964). On Estimating Regression. *Theory of Probability & Its Applications*, 9(1), 141–142. <http://doi.org/10.1137/1109020>

- Paithankar, J. G., Saini, S., Dwivedi, S., Sharma, A., & Chowdhuri, D. K. (2021). Heavy metal associated health hazards: An interplay of oxidative stress and signal transduction. *Chemosphere*, 262, 128350. <http://doi.org/10.1016/j.chemosphere.2020.128350>
- Plackett, R. L., & Hewlett, P. S. (1952). Quantal Responses to Mixtures of Poisons. *Journal of the Royal Statistical Society: Series B (Methodological)*, 14(2), 141–154. <http://doi.org/10.1111/j.2517-6161.1952.tb00108.x>
- R Core Team. (2013). *R: A language and environment for statistical computing: Reference index*. Vienna: R Foundation for Statistical Computing.
- Schulz, E., Speekenbrink, M., & Krause, A. (2018). A tutorial on Gaussian process regression: Modelling, exploring, and exploiting functions. *Journal of Mathematical Psychology*, 85, 1–16. <http://doi.org/10.1016/j.jmp.2018.03.001>
- Siemiatycki, J., & Thomas, D. C. (1981). Biological Models and Statistical Interactions: An Example from Multistage Carcinogenesis. *International Journal of Epidemiology*, 10(4), 383–387. <http://doi.org/10.1093/ije/10.4.383>
- VanderWeele, T. J., & Knol, M. J. (2014). A Tutorial on Interaction. *Epidemiologic Methods*, 3(1), 33–72. <http://doi.org/10.1515/em-2013-0005>
- Wagaman, A. S., & Dobrow, R. P. (2021). *Probability: With Applications and R* (1st ed.). Wiley. <http://doi.org/10.1002/9781119692430>
- Wang, X., Mukherjee, B., & Park, S. K. (2018). Associations of cumulative exposure to heavy metal mixtures with obesity and its comorbidities among U.S. Adults in NHANES 2003–2014. *Environment International*, 121, 683–694. <http://doi.org/10.1016/j.envint.2018.09.035>
- Ward, J. B., Gartner, D. R., Keyes, K. M., Fliss, M. D., McClure, E. S., & Robinson,

W. R. (2019). How do we assess a racial disparity in health? Distribution, interaction, and interpretation in epidemiological studies. *Annals of Epidemiology*, 29, 1–7. <http://doi.org/10.1016/j.annepidem.2018.09.007>

Watson, G. S. (1964). Smooth Regression Analysis. *Sankhyā: The Indian Journal of Statistics*, 26(4), 359–372.